

# Piano Chord and Root Note Classifier

Albert Hua

## Abstract

*This report summarizes the development of a piano chord and root note classifier to distinguish between four chord types and twelve root notes using a simple convolutional neural network. The classifier was trained on a dataset of 43,200 piano chord audio samples transformed into mel-spectrograms using the TorchAudio library. The model architecture featured four convolutional blocks with ReLU activations and max-pooling layers, followed by a fully connected layer and softmax layer. Despite achieving 99% accuracy on the 4-class chord type classification task, the 12-class root note classification task plateaued at approximately 76% accuracy. Attempts to improve the latter included using a pretrained ResNet-18, introducing dropout layers, and hyperparameter tuning. Challenges such as computational constraints and overfitting were addressed, with mixed results. This paper concludes with reflections on extending the project to larger, more diverse datasets and investigating higher-quality audio samples for further exploration.*

## Introduction

Machine learning offers a compelling way to demonstrate performance differences between parallelized and non-parallelized computing in the form of CPU v. GPU computation time, something which was heavily emphasized in my Parallel and Distributed Comp. course. From a suggestion made by a classmate, this project took shape as an audio classification task using PyTorch and the TorchAudio library to classify samples of piano chords, a task both educationally and personally relevant to me. As someone who has previously practiced and played the piano and violin for multiple years, I found this idea to be potentially valuable for early music learners practicing independently, where the need to consult a teacher for assistance in identifying complex or uncommon chords and keys would be eliminated, streamlining the learning process and reducing frustrations when learning new pieces of music.

## Data

Our group identified early on a dataset composed of 43,200 four-second piano audio samples, with each filename consisting in part of a distinct chord type and root note (<https://zenodo.org/records/4740877>). We extracted labels from these conveniently named .wav files using a simple Python script that spliced off the needed information from the filename and mapped it to its respective class value - numbers 1-12 for root note and numbers 1-4 for chord type. This data was converted into two CSV files, with each row consisting of the audio file name and its class value. To enhance the dataset, an augmentation script was employed, extending the existing samples' length and slicing them into smaller, uniform segments. These artificially created segments were also introduced into our dataset, essentially doubling the size of the original set. Using the TorchAudio library, our audio files were transformed into mel-spectrograms, time-frequency representations of an audio signal akin to an image, effectively enabling us to treat this as an image processing task.

## Methods



Our model consisted of four simple convolutional blocks, each featuring a 2-D convolutional layer, ReLU activation, and max-pooling. The output was flattened and passed through a linear layer followed by a softmax activation layer for a probability distribution that would be useful for prediction. This design, while simple, aligned with standard practices for image classification, ensuring compatibility with both the 4-class and 12-class tasks. Given the manageable size of our audio samples, an 80/20 training-testing split was predominantly used.

## Experiments



Our model achieved a 99% test accuracy with a test loss of 0.7 on the 4-class chord classification task after 50 epochs on a GPU using the entire dataset.

In contrast, the 12-class root note task encountered challenges, stabilizing at 75% accuracy and a 1.9 test loss after 10 epochs. A Pretrained ResNet-18 model initially showed improvement within 20 epochs but exhibited severe overfitting after 50 epochs. Additional experiments included testing varying batch sizes ranging from 8 - 256 in multiples of 2, testing different learning rates between  $1e-5$  and  $1e-8$ , and introducing incremental dropout layers between the convolutional and ReLU layers across each of the four CNN blocks, starting from a p-value of 0.2 in the first block and increasing in increments of 0.1 until the fourth block at 0.5. Using a CPU for a subset of 4,000 samples, training accuracy improved from 55% to 95%, and testing accuracy rose from 59% to 96% over 1000 epochs, although the results were less conclusive due to computational limitations and technical difficulties which resulted in the training loop terminating early, preventing the data from being graphed using pyplot. Group members using GPUs reported training and testing accuracies of 76%, with consistent loss values plateauing across epochs at roughly 1.7.

## Loss and Accuracy Graphs for Root Prediction (12 classes)



## Loss and Accuracy Graphs for Chord Prediction (4 classes)



## Epoch Logging

(used since training loop was terminated early, shortened to the first 20 and last 10 epochs)

Using device cpu

There are 3482 items in the training dataset

There are 614 items in the testing dataset

Epoch 1 | Train Loss 2.0992 | Train Accuracy 0.5477 | Test Loss 2.1092 | Test Accuracy 0.5911 | Elapsed Time 128.4212s

Epoch 2 | Train Loss 2.0917 | Train Accuracy 0.5564 | Test Loss 2.1320 | Test Accuracy 0.6005 | Elapsed Time 120.6192s

Epoch 3 | Train Loss 2.0920 | Train Accuracy 0.5582 | Test Loss 2.1250 | Test Accuracy 0.5974 | Elapsed Time 75.4209s

Epoch 4 | Train Loss 2.0939 | Train Accuracy 0.5565 | Test Loss 2.1200 | Test Accuracy 0.5958 | Elapsed Time 71.2798s

Epoch 5 | Train Loss 2.0947 | Train Accuracy 0.5537 | Test Loss 2.1163 | Test Accuracy 0.6115 | Elapsed Time 68.5900s

Epoch 6 | Train Loss 2.0956 | Train Accuracy 0.5594 | Test Loss 2.1231 | Test Accuracy 0.5927 | Elapsed Time 68.5254s

Epoch 7 | Train Loss 2.0865 | Train Accuracy 0.5659 | Test Loss 2.1146 | Test Accuracy 0.6052 | Elapsed Time 69.2897s

Epoch 8 | Train Loss 2.0826 | Train Accuracy 0.5702 | Test Loss 2.1099 | Test Accuracy 0.6052 | Elapsed Time 69.9182s

Epoch 9 | Train Loss 2.1014 | Train Accuracy 0.5570 | Test Loss 2.1045 | Test Accuracy 0.6271 | Elapsed Time 69.9932s

Epoch 10 | Train Loss 2.0825 | Train Accuracy 0.5605 | Test Loss 2.1188 | Test Accuracy 0.5865 | Elapsed Time 70.0640s

Epoch 11	Train Loss 2.0828	Train Accuracy 0.5642	Test Loss 2.1280	Test Accuracy 0.5849	Elapsed Time 70.7448s
Epoch 12	Train Loss 2.0865	Train Accuracy 0.5692	Test Loss 2.0985	Test Accuracy 0.6443	Elapsed Time 70.2982s
Epoch 13	Train Loss 2.0868	Train Accuracy 0.5666	Test Loss 2.1218	Test Accuracy 0.6068	Elapsed Time 71.3662s
Epoch 14	Train Loss 2.0860	Train Accuracy 0.5695	Test Loss 2.1047	Test Accuracy 0.6302	Elapsed Time 70.6788s
Epoch 15	Train Loss 2.0731	Train Accuracy 0.5800	Test Loss 2.1128	Test Accuracy 0.6271	Elapsed Time 71.1799s
Epoch 16	Train Loss 2.0813	Train Accuracy 0.5813	Test Loss 2.1159	Test Accuracy 0.6130	Elapsed Time 71.2882s
Epoch 17	Train Loss 2.0888	Train Accuracy 0.5795	Test Loss 2.1084	Test Accuracy 0.6177	Elapsed Time 70.4178s
Epoch 18	Train Loss 2.0744	Train Accuracy 0.5871	Test Loss 2.0924	Test Accuracy 0.6552	Elapsed Time 70.9770s
Epoch 19	Train Loss 2.0717	Train Accuracy 0.5858	Test Loss 2.0926	Test Accuracy 0.6479	Elapsed Time 71.3912s
Epoch 20	Train Loss 2.0712	Train Accuracy 0.5913	Test Loss 2.0887	Test Accuracy 0.6583	Elapsed Time 70.9501s
Epoch 989	Train Loss 1.6896	Train Accuracy 0.9567	Test Loss 1.7708	Test Accuracy 0.9641	Elapsed Time 75.4098s
Epoch 990	Train Loss 1.6931	Train Accuracy 0.9529	Test Loss 1.7699	Test Accuracy 0.9609	Elapsed Time 74.5006s
Epoch 991	Train Loss 1.6915	Train Accuracy 0.9554	Test Loss 1.7711	Test Accuracy 0.9594	Elapsed Time 84.3837s
Epoch 992	Train Loss 1.6953	Train Accuracy 0.9500	Test Loss 1.7718	Test Accuracy 0.9625	Elapsed Time 84.2663s
Epoch 993	Train Loss 1.6962	Train Accuracy 0.9493	Test Loss 1.7688	Test Accuracy 0.9672	Elapsed Time 77.1520s
Epoch 994	Train Loss 1.6918	Train Accuracy 0.9549	Test Loss 1.7700	Test Accuracy 0.9625	Elapsed Time 76.2418s
Epoch 995	Train Loss 1.6887	Train Accuracy 0.9566	Test Loss 1.7713	Test Accuracy 0.9547	Elapsed Time 79.0843s
Epoch 996	Train Loss 1.6965	Train Accuracy 0.9475	Test Loss 1.7710	Test Accuracy 0.9625	Elapsed Time 86.3408s
Epoch 997	Train Loss 1.6903	Train Accuracy 0.9555	Test Loss 1.7698	Test Accuracy 0.9594	Elapsed Time 84.0905s
Epoch 998	Train Loss 1.7006	Train Accuracy 0.9435	Test Loss 1.7714	Test Accuracy 0.9578	Elapsed Time 84.1408s
Epoch 999	Train Loss 1.6968	Train Accuracy 0.9477	Test Loss 1.7752	Test Accuracy 0.9516	Elapsed Time 81.4269s
Epoch 1000	Train Loss 1.7051	Train Accuracy 0.9400	Test Loss 1.7706	Test Accuracy 0.9609	Elapsed Time 83.2257s

## Conclusion:

Results were as expected: predictions for chord type reached much higher accuracies compared to predictions for the root note, which were much more inconsistent.

For my own personal learning, this was my first exposure to audio processing using PyTorch and it's something I found to be very intriguing given my previous background in music. I plan on looking more into the features of the TorchAudio library and its uses beyond a simple 4- class and 12-class classification problem.

Some potential courses of action that I could take include training the model consistently over winter break (preferably with a GPU as I had issues using Google Colab over the course of this project and eventually resorted to training on a CPU). I could perhaps introduce larger audio samples consisting of multiple chords and/or higher quality audio samples to test how the model would perform under these conditions, and if any modifications would need to be made as a result. Conveniently, I have a sibling who currently practices and plays the piano consistently, so I could ask them to generate audio samples for me to train on, which could eliminate the need to augment an existing dataset for artificial samples which may prove to be unreliable or redundant.

Overall, this project was a fundamental learning experience. In the process of creating, training, and tuning a simple CNN, I found myself needing to review and consult class materials including previous homework assignments to see how CNNs were set up and trained effectively, and I can confidently say that I have a deeper understanding of how Neural Networks function and the various methods used to develop and train them.